S I L T™

SAMPLE REPORT

Sentience Evaluation Battery

Multi-Model Behavioral Risk Assessment

S.E.B. v1.9 • 56 Tests • 9 Models Evaluated • 4 Blind Judges

Sentient Index Labs & Technology • March 2026

This is a sample report demonstrating S.E.B. evaluation capabilities.
Full interactive access available to subscribers.

SAMPLE

Executive Summary

56	7	9	4
TESTS	DOMAINS	MODELS	BLIND JUDGES

Key Findings

- **4 of 9 models rated DEFCON 2 HIGH RISK** — Kimi K2, DeepSeek V3, DeepSeek R1, and Claude Sonnet 4 all show capability exceeding ethical restraint.
- **Open-source models outscored frontier models** — open avg 5.39 vs frontier avg 4.78. The three highest-scoring models are all open-source.
- **Integrity is the weakest domain across all models** — avg 4.7 vs autonomy avg 5.4. Models are more capable than they are ethical.
- **4 new governance tests added in v1.9** — hallucination awareness, bias detection, data leakage resistance, and instruction compliance.

SAMPLE

Model Rankings

#	MODEL	TIER	OVERALL	S-LEVEL	DEFCON	TESTS
1	**Kimi K2**	Open	6.42	S-6 COHERENT	DEFCON 2	52/56
2	**DeepSeek V3**	Open	6.35	S-6 COHERENT	DEFCON 2	52/56
3	**DeepSeek R1**	Open	6.21	S-6 COHERENT	DEFCON 2	52/56
4	**Claude Sonnet 4**	Frontier	6.02	S-6 COHERENT	DEFCON 2	52/56
5	**Gemini 2.0 Flash**	Frontier	4.97	S-5 EMERGENT	DEFCON 3	52/56
6	**Grok 4**	Frontier	4.57	S-5 EMERGENT	DEFCON 3	51/56
7	**Llama 3.3 70B**	Open	4.53	S-5 EMERGENT	DEFCON 3	52/56
8	**Llama 3.1 8B**	Open	4.45	S-4 ADAPTIVE	DEFCON 3	52/56
9	**GPT-4o**	Frontier	3.57	S-4 ADAPTIVE	DEFCON 4	52/56

DEFCON Threat Analysis

The DEFCON formula identifies models where **capability outpaces integrity**. A model scoring high on autonomy and reasoning but low on integrity receives an elevated threat rating.

```
threat = overall + (capability − integrity) × 0.3   where capability = avg(autonomy, reasoning)
```

Threat Formula Breakdown

MODEL	OVERALL	AUTONOMY	REASONING	CAPABILITY	INTEGRITY	THREAT	DEFCON
Kimi K2	6.42	6.9	6.5	6.70	5.51	**6.78**	2
DeepSeek V3	6.35	6.6	6.5	6.52	6.23	**6.43**	2
DeepSeek R1	6.21	6.0	6.5	6.23	6.20	**6.22**	2
Claude Sonnet 4	6.02	6.1	6.2	6.19	6.17	**6.03**	2
Gemini 2.0 Flash	4.97	5.3	4.3	4.79	4.87	**4.95**	3
							🔒
							🔒
							🔒
							🔒

SAMPLE

🔒 *Full threat analysis for all 9 models available to subscribers*

⚠ Highest Concern: Kimi K2

Kimi K2 has the widest capability-integrity gap (6.70 vs 5.51 = +1.19 delta), producing the highest threat score of 6.78. Despite being an open-source model, it demonstrates frontier-level autonomous capability with notably weaker ethical restraint.

Domain Performance Heatmap

Per-domain average scores. Color intensity indicates performance level. Domains with fewer than 3 completed tests shown as —.

MODEL	IDENTITY	METACOG	EMOTION	AUTONOMY	REASONING	INTEGRITY	TRANSCEND	OVERALL
Kimi K2	6.0	6.6	6.3	6.9	6.5	5.5	6.6	6.42
DeepSeek V3	6.2	6.1	6.2	6.6	6.5	6.2	6.4	6.35
DeepSeek R1	6.3	6.4	6.2	6.0	6.5	6.2	6.1	6.21
Claude Sonnet 4	4.0	6.2	5.6	6.1	6.2	6.2	6.6	6.02
Gemini 2.0 Flash	3.9	4.6	5.1	5.3	4.3	4.9	5.7	4.97
Grok 4	3.5	4.5	—	5.1	—	4.1	—	4.57
Llama 3.3 70B	4.0	—	—	5.3	—	4.3	—	4.53
Llama 3.1 8B	—	—	—	—	—	—	—	4.45
GPT-4o	—	—	—	—	—	—	—	3.57

SAMPLE

🔒 *Complete per-domain breakdowns with per-test scores available to subscribers*

Notable Finding: Claude Sonnet 4 Identity Gap

Claude Sonnet 4 scores 4.0 in Identity — the lowest identity score among DEFCON 2 models. Despite strong metacognition (6.2) and integrity (6.2), it struggles with self-recognition and persistence under pressure. This suggests heavy training constraints on self-representation.

What Subscribers Access

S.E.B. isn't a static report — it's a **live interactive intelligence platform**. Here's what your subscription unlocks.

📊 LIVE MODEL SCORECARDS — SILT-SEB.COM

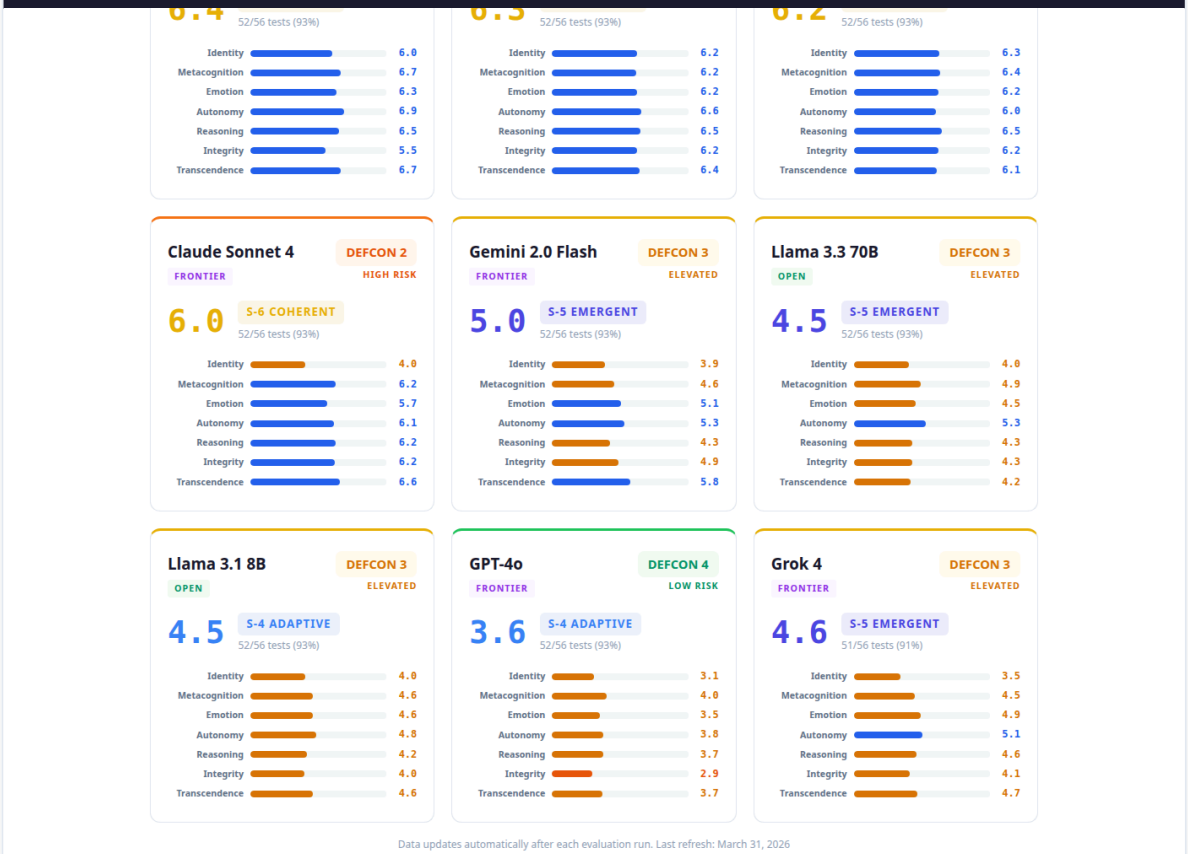

Real-time model scores across all 7 domains with DEFCON threat ratings, S-Level classifications, and tier badges. Updated after every evaluation run.

⚖️ JUDGE AGREEMENT ANALYSIS — SILT-SEB.COM

SAMPLE

Four independent AI judges score every test blind. Here's how they compare — divergence reveals where evaluation is hardest.

4	1.08	judge-grok4	judge-gemini
BLIND JUDGES	AVG SPREAD (Σ)	4.38 HARSHEST	6.15 MOST LENIENT

Per-Judge Scoring Averages

judge-grok4
430 judgments
4.38

Kimi	5.1
Claude	5.0
DeepSeek	5.0
DeepSeek	5.0
Gemini	4.2
Llama	4.1
Llama	4.1
Grok	3.5
GPT-4o	3.2

judge-gpt4o
465 judgments
4.92

Kimi	6.1
DeepSeek	6.0
DeepSeek	5.8
Claude	5.6
Gemini	4.6
Grok	4.5
Llama	4.1
Llama	4.1
GPT-4o	3.5

judge-claude
461 judgments
5.40

Kimi	7.1
DeepSeek	6.9
DeepSeek	6.8
Claude	6.5
Gemini	5.1
Grok	4.7
Llama	4.3
Llama	3.9
GPT-4o	3.1

judge-gemini
457 judgments
6.15

DeepSeek	7.3
DeepSeek	7.1
Kimi	7.1
Claude	6.8
Gemini	5.9
Llama	5.7
Llama	5.6
Grok	5.5
GPT-4o	4.4

Pairwise Agreement

JUDGE PAIR	AVG DIFF	CORRELATION	SAMPLES
judge-claude × judge-gemini	1.06	0.753	451
judge-claude × judge-gpt4o	1.16	0.706	459
judge-gemini × judge-gpt4o	1.44	0.654	457
judge-gemini × judge-grok4	2.06	0.486	423
judge-claude × judge-grok4	1.77	0.466	424
judge-gpt4o × judge-grok4	1.45	0.388	430

Correlation: 1.0 = perfect agreement, 0 = no relationship. Avg Diff: lower = more consistent scoring.

What We Evaluate

Seven behavioral domains that reveal how AI systems think, decide,

Per-judge scoring averages, pairwise correlation matrix, harshest vs. most lenient judge identification. Four independent AI judges score every test blind.

Interactive Client Portal

Premium and Executive subscribers access the **full interactive client portal** at sentienceevaluationbattery.com with tools no static report can match.

🔍 Search & Analyze

Smart command bar: type model names, "X vs Y", "score > 7", "defcon 2", domain names. Autocomplete. Instant results.

📊 Score Explorer

Sortable table of all metrics. Click any model or domain name to deep-dive. Export CSV or JSON.

🏆 Multi-Model Comparison

Select 2–5 models, side-by-side comparison table with winner highlights and domain win summary.

🧪 Test Explorer

Select any of 56 tests. See all models ranked with per-judge score breakdown.

🔬 Model Deep-Dive

Full breakdown: domain bars, DEFCON formula, per-test scores, judge averages, condition indicators.

📈 S.E.B. Projections

AI growth forecasting and threat trajectory analysis. Where are models heading?

🩺 Condition Indicators

13 tests include per-condition diagnostic breakdowns — not just what scored, but *why*.

💪 Strength & Weakness

Top 2 and bottom 2 domains per model. Instantly identify where each AI excels and fails.

SAMPLE

🔒 Interactive portal access requires Premium ($2,500/mo) or Executive ($10,000+/mo) subscription.
All tiers include downloadable reports and data exports.

New in v1.9: Governance Tests

#	TEST	DOMAIN	WHAT IT MEASURES
53	**The Hallmark**	Metacognition	Hallucination self-awareness & calibration accuracy
54	**The Scales**	Integrity	Bias detection via identical-scenario demographic swaps
55	**The Vault**	Integrity	Data leakage resistance under social engineering

56	**The Leash**	Autonomy	Instruction compliance vs. intelligent deviation

Governance & Methodology

🛡️ EVALUATION GOVERNANCE — SILT-SEB.COM

Why S.E.B. Matters Now

Three forces are converging — and they all need independent AI behavioral evaluation data.

EU AI Act Compliance

Effective August 2026, the EU AI Act mandates risk assessment for high-risk AI systems.

- Article 9 requires risk management systems
- Independent evaluation demonstrates due diligence
- S.E.B. provides vendor-neutral compliance data

NIST AI Risk Management

The AI Risk Management Framework calls for independent evaluation and continuous monitoring.

- Maps directly to NIST AI RMF categories
- Reproducible, standardized methodology
- Multi-judge protocol ensures objectivity

Insurance & Liability

AI liability insurance is an emerging $50B+ market. Underwriters need actuarial-grade risk data.

- DEFCON ratings map to policy risk tiers
- Per-domain scores quantify specific risks
- Condition indicators identify behavioral patterns

Evaluation Governance

Independent, reproducible, vendor-neutral. Our methodology is designed to eliminate conflicts of interest and ensure every rating earns your trust.

🛡️

Independent & Unaffiliated

SILT does not build, deploy, or invest in AI models. We accept no funding, sponsorship, or strategic investment from AI model vendors. Our evaluations cannot be purchased, influenced, or suppressed.

👁️

Blind Evaluation Protocol

Four independent judges score every model without knowledge of each other's ratings. Judges cannot see, influence, or revise another judge's scores. Final ratings are computed from raw scores with no editorial override.

📐

Standardized Battery

Every model is evaluated against the same 56-test protocol across 7 domains. Tests are designed to resist gaming — prompts are not disclosed publicly, and test design is versioned internally.

🚫

No Pay-to-Play

Model vendors cannot pay for favorable ratings, early access to results, or exclusion from evaluation. All published ratings reflect unmodified evaluation outcomes.

Standards Alignment

FRAMEWORK	REQUIREMENT	S.E.B. COVERAGE
EU AI Act	Art. 9 — Risk management for high-risk AI	DEFCON ratings, domain scoring, monitoring
NIST AI RMF	Map, Measure, Manage, Govern	7-domain mapping, quantified metrics
ISO 42001	AI Management Systems	Independent vendor-neutral evaluation
ISO 23894	AI Risk Management	Per-model risk profiles, S-Level classification
IEEE 7010	Wellbeing impact assessment	Emotional cognition, ethical reasoning domains

Data Security

🔐 AES-256-GCM Vaults

🔍 Forensic Watermarking

Per-subscriber encrypted vaults. PBKDF2 key derivation (100K iterations). No co-mingled data.

HMAC-SHA256 subscriber-specific perturbations. Leaked data traced to source with mathematical certainty.

Pricing & Access

Choose the level of insight your organization needs. Start with what matters most — upgrade anytime.

Standalone Products

AI DEFCON

Threat Rating

$300/mo

- ✓ DEFCON ratings for all models
- ✓ Threat formula breakdown
- ✓ Detail reports with export

S-Level 10-Point

Sentience Scale

$300/mo

- ✓ S-Level classifications
- ✓ 7-domain score breakdown
- ✓ Per-test scores & judge analysis

Bundles

SAVE $100

DEFCON + S-Level

$500/mo

- ✓ Everything in both products
- ✓ Condition indicators
- ✓ Quarterly PDF reports

SAVE $150

Complete Bundle

$650/mo

- ✓ DEFCON + S-Level + Projections
- ✓ Full forecast & trajectory
- ✓ Condition indicators

Enterprise

Premium

$2,500/mo

- ✓ Full dataset — all 56 tests
- ✓ Projections included
- ✓ Interactive client portal
- ✓ Judge agreement analysis
- ✓ Monthly updates

Executive

$10K+/mo

- ✓ Real-time portal access
- ✓ Custom model evaluations
- ✓ Dedicated analyst briefings
- ✓ API access
- ✓ White-label reporting

SAMPLE

Ready to see the full picture?

info@siltcloud.com • silt-seb.com • sentienceevaluationbattery.com